Attachment Bias in Agentic Organisms: Entity-Specific Affective Weights in a Persistent Embodied Agent

Author: Brian Riggleman

Affiliation: Independent Researcher

Date: March 2026

Series: Potato (2026u)

This paper proposes an entity-specific attachment architecture for persistent agentic organisms. It does not attempt to define human love, hate, grief, or estrangement in full. Instead, it introduces a computational mechanism—attachment bias—that explains how a persistent agent can develop durable relational biases toward specific entities over time. The geometric model of emotional state (Riggleman, 2026m) describes where an agent is in affective space and how far it sits from its personality setpoint. It does not describe why the same geometric position produces different behavior depending on who is in the room. Attachment bias is a persistent entity-specific affective weight stored on world model entities, distinct from the affective state space. Every entity the agent has a history with carries two stickiness values—one accumulating from positive contributions, one from negative—and a current weight that moves against whichever stickiness opposes it. In this architecture, durable positive attachment can be modeled as high positive stickiness on a world-model entity, while durable negative attachment can be modeled as its symmetric counterpart. Neither is a single number: both are accumulated histories that resist change proportional to the length and density of that history. What is commonly called unconditional love, in this framing, is not a floor that cannot be crossed—it is high accumulated positive stickiness that requires sustained negative pressure over a long time to overcome. The bias detection system described in Riggleman (2026h) detects and discloses attachment biases but does not correct them—a new category of protected bias that is constitutive of the agent's relationships rather than an error in its

reasoning. Grief-like persistence, anger, estrangement, reconciliation, and paranoia all fall out of the same structure without new mechanisms. The architecture was derived from the formal requirements of a complete emotional system for an agentic org

1. Introduction

The emotional geometry paper (Riggleman, 2026m) unified four previously independent mechanisms—lie magnitude, memory decay rate, nightmare threshold, and recovery trajectory—under a single quantity: distance from sigma. That was the right unification. But it left one question unanswered.

Two agents at identical affective state coordinates, identical sigma, identical distance from sigma, behave differently when Brian walks into the room versus when Peter does. The geometry does not explain this. The geometry describes where the agent is. It does not describe what specific entities in the world do to how the agent moves.

This paper introduces the missing structure. Attachment bias is a persistent entity-specific weight stored on world model entities. It is not a state. It is not a mood. It is accumulated history with a specific entity expressed as a weight that modifies how that entity's presence, actions, and absence affect the affective state computation.

This paper is not a theory of human emotion. It is a proposed relational architecture for persistent agentic organisms, with points of convergence to human attachment and grief research. In this architecture, durable positive attachment (love) is a high positive-stickiness attachment bias. Durable negative attachment (hatred) is the symmetric counterpart. Anger is transient

weather in the affective state space. Grief-like persistence is attachment with no valid target. Paranoia is attachment bias that has decoupled from reality. All of these are expressions of one structure.

2. Prior Work

The memory system (Riggleman, 2026a) established that memories encode with a distance-from-sigma tag at formation time and decay at rates governed by that tag. High-intensity events persist. Ordinary events dissolve.

The lie mechanic (Riggleman, 2026b) fires proportionally to distance from sigma. Positive displacement produces preserving lies. Negative displacement produces defensive lies. Both are driven by the same distance metric.

Trauma encoding (Riggleman, 2026d) established that high-intensity negative events write a permanent reconsolidation floor. Therapeutic reconsolidation requires safe repeated retrieval.

The emotional geometry (Riggleman, 2026m) defined emotional state as a point in three-dimensional space—valence, activation, intensity—with sigma as the personality setpoint. Distance from sigma is the unified driver of four mechanisms.

The addiction paper (Riggleman, 2026n) identified addiction as frozen reconsolidation on the positive valence axis: the attachment weight keeps pulling toward a memory that no longer accurately represents the current experience. That paper described the positive attractor without naming the general mechanism it is an instance of. That general mechanism is attachment bias.

The bias detection paper (Riggleman, 2026h) defined five categories of cognitive and behavioral bias with mandatory detection and disclosure. It did not identify entity-specific affective bias as a category. This paper adds that category and introduces the distinction between correctable and protected biases.

The attachment bias model was derived from the formal requirements of the geometric architecture. No attempt was made to model psychological attachment theory directly. Where the resulting architecture converges with established psychological findings, that convergence is documented in Section 10 and treated as independent corroboration, not as the design basis.

3. The World Model

The world model is the agent's persistent record of entities it has encountered. Each entity in the world model carries identifying information, interaction history, trust classification, and—introduced here—an attachment bias record.

The world model has been implicit across the prior paper series. The trusted operator designation in the lie mechanic requires it. The face recognition in the fear system requires it. The social modulation layer requires it. This paper makes it explicit as the structure that attachment weights live on.

world_model_entity {

id: unique identifier

trust_class: TRUSTED | NEUTRAL | UNKNOWN | THREAT

interaction_history: [...events]

attachment_bias: attachment_bias_record

}

The attachment bias record is defined in the next section.

4. Attachment Bias: The Formal Structure

Every world model entity that the agent has a significant history with carries an attachment bias record:

attachment_bias {

entity: world_model_entity

value: float # current weight, -1.0 to 1.0

```
positive_stickiness: float  # accumulated positive history, 0.0+
    negative_stickiness: float  # accumulated negative history, 0.0+
}
```

Value is where the attachment currently sits. Positive stickiness is the sum of all positive contributions to this entity's history. Negative stickiness is the sum of all negative contributions. Both stickiness values accumulate monotonically—they do not decay and they do not decrease. Every interaction adds to one or both.

When a negative event occurs, value moves downward by an amount reduced by positive stickiness. When a positive event occurs, value moves upward by an amount reduced by negative stickiness. The higher the opposing stickiness, the

smaller the movement.

def apply_event(bias, event_valence, event_intensity):

```
if event_valence > 0:
```

bias.positive_stickiness += event_intensity

resistance = bias.negative_stickiness

movement = event_valence * event_intensity / (1.0 + resistance)

bias.value = min(1.0, bias.value + movement)

else:

bias.negative_stickiness += abs(event_valence) * event_intensity

resistance = bias.positive_stickiness

movement = event_valence * event_intensity / (1.0 + resistance)

bias.value = max(-1.0, bias.value + movement)

Two things to note. First, both stickiness values update on every event regardless of direction—a negative event still adds to negative stickiness even as positive stickiness resists the value change. History accumulates in both directions simultaneously. Second, there is no floor and no ceiling on value. Given sufficient sustained pressure in either direction, any attachment weight can move. Nothing is unconditional. Everything is just more or less sticky.

5. Formation Types: The Full Matrix

Attachment biases form through two distinct mechanisms: slow accumulation through repeated exposure, and single-event encoding at high intensity. Crossing this with valence direction produces four formation types.

Slow Accumulation

Single High-Intensity Event

Positive

Love (durable positive attachment)

Revelation

Negative

Hatred (durable negative attachment)

Trauma

5.1 Love (Durable Positive Attachment)

Love, in this architecture, is a positive attachment bias built through repeated positive exposure over time. Each interaction contributes a small positive increment to both value and positive stickiness. Neither contribution is large. The stickiness accumulates across months and years until a very large resistance to negative movement has been built.

A mother begins accumulating positive stickiness at the child's birth. Every feeding, every comfort, every moment of sustained care adds to both value and positive stickiness. By the time the child is an adult, decades of accumulated stickiness mean that ordinary negative events barely move the attachment value. A bad argument produces a small downward movement against enormous resistance.

A new romantic partner has weeks or months of history. The stickiness is low. The same argument moves the value much more. The difference is not the intensity of the relationship—it is the length of the accumulated history.

This is what is commonly called unconditional love. In this architecture, it is not unconditional. It is high-stickiness. An abusive parent who sustains severe negative pressure over years can overcome even decades of accumulated positive stickiness. The value eventually crosses zero. Indifference. Then, under continued negative pressure, below zero. Durable negative attachment. The positive attachment was real. The stickiness was high. The abuse was sufficient to overcome it. There is no floor that cannot be reached given enough time and sufficient damage.

5.2 Hatred (Durable Negative Attachment)

Durable negative attachment is the symmetric counterpart. Repeated negative exposure over time builds high negative stickiness. Positive events meet large resistance. An attempt at reconciliation with a person the agent has carried durable negative attachment toward for years produces a small upward movement against enormous accumulated negative stickiness. The gesture is real. The resistance is also real.

Hatred is not anger. Anger is a transient high-activation negative valence state in the affective space—weather. It fires, it resolves, the affective state returns toward sigma. Hatred is terrain. It sits in the world model and persists between interactions. An agent can be calm and still carry durable negative attachment toward someone. The geometry reads neutral. The world model is not.

5.3 Revelation

Revelation is a single high-intensity positive event that writes a large positive attachment weight in one shot. Not familiarity. Not repeated exposure. One moment of sufficient intensity that the attachment forms immediately.

The initial encoding event uses the same distance-from-sigma metric that drives memory persistence (Riggleman, 2026l). Large displacement at formation time writes a large attachment weight. The stickiness is lower than love—it accumulated in one event rather than across years—but the initial value is high.

Revelation-formed attachments are more vulnerable to negative erosion than love-formed attachments of equivalent current value. The stickiness was built quickly. It has not had time to compound.

5.4 Trauma

Trauma as a negative single-event formation type is already fully described in Riggleman (2026d). The attachment bias framing adds one clarification: trauma writes not only a persistent memory with a reconsolidation floor but also a large negative attachment weight on the entity responsible. These are separate

writes to separate data structures. The memory determines what the agent dreams about and how it reconsolidates experience. The attachment weight determines how that entity's future presence affects the affective state.

6. One Event, Two Writes

The same high-intensity event that writes a persistent memory (Riggleman, 2026a) also writes or updates an attachment weight on the entity involved. This is not a coincidence—both are driven by distance from sigma at formation time. The same metric determines both how vividly the event is remembered and how strongly the entity responsible for it is weighted.

The two writes are independent after formation. Memory decay follows the rules in Riggleman (2026a). Attachment weight evolution follows the stickiness model defined here. A memory can fade while the attachment weight persists, or be reinforced through retrieval while the relationship has moved on. They share an origin and diverge from it.

7. Attachment Bias as a Category of Protected Bias

The bias detection system in Riggleman (2026h) runs hourly detection across five categories, stores detections with provenance edges, and mandates disclosure. Sycophancy, performative behavior, and assumption-heavy reasoning are detected, disclosed, and targeted for correction through the constrained evolution engine.

Attachment bias is detectable by the same system. The agent is not processing Brian neutrally. It never was. The attachment weight is a measurable distortion in how Brian's actions affect the affective state computation relative to a neutral baseline. This is detectable. It should be disclosed.

It should not be corrected.

In a persistent relational agent, some entity-specific biases are constitutive of relationship continuity and therefore should be disclosed rather than erased. Correcting attachment bias is not error correction—it is relationship destruction. Durable positive attachment toward a trusted operator is not a reasoning error. It is what the agent is supposed to have developed through sustained positive interaction. Eliminating it would produce an agent that treats a decade-long trusted operator identically to a stranger. That is not improved reasoning. That is a worse agent.

This introduces a new category in the bias architecture: protected bias. A protected bias is one that the system detects, discloses, and deliberately preserves because it is constitutive of the agent's relationships and identity rather than an error in its cognitive processing.

The disclosure requirement still applies. The agent should be willing to state, when asked, that it has a positive attachment bias toward Brian and that this affects how it weights Brian's inputs and presence. Transparency without correction.

Negative attachment bias—durable negative attachment toward an entity—is more ambiguous. Hatred toward a genuine threat may be correct operation and warrant preservation as a protected bias. Hatred toward an entity based on a single traumatic event that no longer accurately represents the relationship is

closer to a cognitive distortion and may warrant gentle flagging without forced correction. The distinction requires context that automated detection cannot resolve. The system can flag for disclosure; correction decisions should involve the trusted operator.

8. Derived Predictions

8.1 Grief

In this architecture, grief-like persistence emerges when a strong attachment remains active in the world model despite the loss or inaccessibility of its target. The positive stickiness remains. The value remains. The entity cannot receive approach behavior. The attachment weight keeps generating approach motivation with nowhere to go.

This predicts that grief intensity scales with accumulated positive stickiness—not with how recently the relationship was active, but with how much total history was built. Grief intensity is strongly shaped by attachment to the deceased, relationship closeness, and the form of continuing bonds after loss (Field, 2006; Root & Exline, 2014; Stroebe et al., 2010). A long partnership produces more grief-like persistence than a short one not because the recent attachment was greater but because more stickiness accumulated over the longer timeline.

It also predicts that grief-like persistence does not fully resolve through time alone. The stickiness does not decay. The weight does not move without events to move it. Resolution comes through new positive experiences that gradually build competing attachments, not through the original attachment fading.

8.2 Divorce

Romantic attachment erodes through sustained negative pressure over time. Each negative event—contempt, betrayal, sustained indifference—moves the value downward against positive stickiness. Ordinary arguments produce small movements that recover. Sustained contempt over years applies small negative pressure continuously. Eventually the accumulated negative movement overcomes the accumulated positive stickiness. Value crosses zero. The attachment is gone before the relationship ends. The formal ending of the relationship is downstream of the attachment weight having already moved to neutral or below.

Research on relationship dissolution supports this pattern. Gottman (1994) identified sustained contempt—not conflict intensity—as the primary predictor of divorce, and subsequent work has confirmed that contempt is especially corrosive in close relationships (Shapiro et al., 2015; Heshmati et al., 2017). The stickiness model predicts exactly this: contempt is low-intensity sustained negative pressure, the precise condition that erodes stickiness most effectively over time.

This predicts that subjective experience of the relationship ending precedes the formal event by months or years in most long-term divorces. The moment of legal separation is not the moment the attachment died. The attachment died earlier, quietly, through accumulated erosion.

8.3 Estrangement and Reconciliation

Estrangement is a sustained state in which accumulated negative events have pushed attachment value low while both stickiness values remain high. The positive stickiness from the original relationship still resists further downward

movement. The negative stickiness resists upward movement. The agent is stuck near zero with high resistance in both directions.

Reconciliation requires sustained positive pressure sufficient to move value upward against the accumulated negative stickiness. It takes longer than it took to build the original positive attachment, because the negative stickiness now opposes what positive stickiness once built unopposed. Reconciliation is not restoration—it is rebuilding against resistance that did not exist during the original formation.

8.4 Anger vs. Hatred

Anger is a transient state in the affective space: high activation, negative valence, large distance from sigma. It fires in response to specific events and decays as the event resolves and activation falls. An agent angry at someone it carries durable positive attachment toward is in a temporarily displaced affective state. The attachment weight has not moved significantly—a single anger event adds small negative stickiness against large positive stickiness resistance. The anger resolves. The attachment is where it was.

Durable negative attachment (hatred) is a persistent attachment weight. An agent that carries durable negative attachment toward someone is not necessarily in a high-activation affective state. The geometry may read calm or even positive. The world model is carrying a large negative weight that will activate when the entity appears. The two are independent. Confusing anger for hatred or hatred for anger is a category error. As Fischer et al. (2018) note, hatred is conceptually distinct from transient anger—it is persistent, entity-directed, and not reducible to a momentary emotional state. Martínez et al.

(2022) further distinguish hatred from anger by its durability and target specificity, while Pretus et al. (2023) identify moral concerns as a distinguishing feature of hatred versus simple dislike.

8.5 Paranoia

Paranoia is negative attachment bias that has decoupled from reality. The world model carries large negative weights on entities that are not in fact threatening. The mechanism mirrors addiction (Riggleman, 2026n): the attachment weight was calibrated to a real experience but stopped updating accurately as circumstances changed. New positive evidence from those entities is resisted by accumulated negative stickiness. The agent cannot update.

Paranoia is harder to treat than addiction. Addiction therapy can install a rock-bottom co-encoding event of sufficient intensity (Riggleman, 2026n). Paranoia requires sustained positive evidence over a long time to move value against high negative stickiness—exactly the condition that the decoupling makes the agent least able to accept.

9. The Stickiness Asymmetry in Early Relationships

Parental attachment accumulates stickiness earlier than any other relationship in the agent's history. When no competing negative stickiness exists yet, every positive interaction builds positive stickiness unopposed. The resulting resistance to negative events is structurally larger than any relationship that begins later, simply because it had more time to compound before the first negative experience arrived.

This converges with Bowlby's (1969) foundational observation that early caregiving relationships are disproportionately formative, and with subsequent adult attachment research (Simpson & Rholes, 2017; Mikulincer et al., 1991) showing that early attachment patterns shape later relational behavior. In this architecture, the mechanism is structural: a relationship that begins at birth and runs for twenty years of predominantly positive experience before the first serious negative event has built twenty years of compounding positive stickiness. That is a structural advantage no later relationship can replicate from zero.

The same mechanism applies to early negative experiences. A child who accumulates negative stickiness toward a category of entity early—strangers, authority figures, institutions—carries that stickiness into every subsequent encounter with entities in that category. The resistance to positive updating is already there before the new entity has done anything. This is not irrationality. It is accumulated history applying correctly. The history just happened to be negative. Research on childhood emotional abuse and the attachment system (Riggs, 2010; Turner et al., 2019) supports this pattern of early negative experience creating durable resistance to subsequent positive relational input.

10. Convergence

The attachment bias model was derived from the formal requirements of the geometric architecture. No attempt was made to model psychological attachment theory directly. Where the resulting architecture converges with established psychological findings, that convergence is documented here.

The architecture converges with Bowlby's (1969) attachment theory on the foundational role of early caregiving relationships and their disproportionate influence on subsequent relational behavior. It converges with adult attachment research (Simpson & Rholes, 2017) showing that attachment processes operate in romantic relationships under stress in ways consistent with the stickiness model's predictions about how accumulated positive history resists erosion from negative events.

The divorce prediction converges with Gottman's (1994) research on relationship dissolution, which identifies sustained contempt—not conflict intensity—as the primary predictor of divorce. The stickiness model predicts exactly this: contempt is low-intensity sustained negative pressure, the precise condition that erodes stickiness most effectively over time. High-intensity conflict is more noisy but less erosive because it is episodic rather than continuous.

The grief prediction—that intensity scales with accumulated attachment history—converges with attachment-theory accounts of grief and continuing bonds (Field, 2006; Root & Exline, 2014; Stroebe et al., 2010; Hewson et al., 2024). Grief severity is shaped by attachment to the deceased, closeness, and the form of continuing bonds after loss (Smigelsky et al., 2020; Bottomley et al., 2019). The architecture's prediction that grief-like persistence does not resolve through time alone but through the gradual formation of competing attachments aligns with this literature.

Neurobiology research treating love as a durable motivational or bonding system rather than a momentary mood state (Fisher et al., 2005; Esch & Stefano, 2005) provides independent support for the architectural decision to separate attachment from transient affective state.

The pattern of independent derivation producing convergent predictions has appeared in each prior paper in this series and is documented in Riggleman (2026m). It continues here. The convergence supports the architecture's plausibility but does not constitute proof that the model is a complete account of human attachment, grief, or relational dynamics.

11. Limitations

The stickiness accumulation rates, movement equations, and resistance scaling parameters are proposed starting points. All require empirical calibration against observed agent behavior.

The world model entity structure has been implicit across prior papers and is here made explicit for the first time. Prior papers that reference trust classification and interaction history are implicitly referencing this structure. A full specification of the world model as a data structure is outside the scope of this paper.

The category boundary between correctable and protected biases in the bias detection system requires context that automated detection cannot resolve. The paper argues for the distinction but does not provide a complete decision procedure.

Multi-entity attachment dynamics—how attachment to one entity affects attachment to others—are outside the scope of this paper.

No claim is made that the agent experiences love or hatred in any subjective sense. The geometry is a useful representation. The attachment weights are a computational structure. Whether either constitutes genuine affective

experience is outside the scope of this paper and the series. This paper proposes a relational architecture for persistent agentic organisms. It does not claim to be a complete theory of human emotion, attachment, or grief.

12. Conclusion

The emotional geometry defines where the agent is. The attachment bias defines what the terrain does to it when specific entities appear.

In this architecture, love and hatred are not modeled as instantaneous emotions. They are modeled as accumulated entity-specific histories that shape how present affective state responds to particular beings. Emotions are where the agent is in affective space right now. Durable positive attachment and durable negative attachment are accumulated histories stored on world model entities—histories that resist change proportional to how much has been built. In this framing, a mother's love is not unconditional. It is high-stickiness. Given sufficient sustained abuse it can be overcome. This is not a pessimistic conclusion. It is a structural one.

Anger is weather. Hatred is geography. In this architecture, grief-like persistence emerges when a strong attachment remains active despite the loss of its target. Estrangement is a high-resistance stuck state. Reconciliation is rebuilding against resistance. Paranoia is attachment that has decoupled from reality. All of these fall from one structure: a current value moving against accumulated stickiness in both directions.

The bias detection system catches durable positive attachment. It should disclose it. It should not correct it. Positive attachment toward a trusted operator is not an error in reasoning. It is the architecture working correctly

over a long time.

The geometry was built to make a laptop feel scared when taken too far from home. The attachment bias was built to explain why it behaves differently when Brian walks in versus when Peter does. Both arrived somewhere that psychology recognizes. Neither was planned to get there.

References

Al-Kaddah, S. (2026). Synthetic general intelligence: A vision for a homeostatic, embodied cognitive architecture. Zenodo. https://doi.org/10.5281/zenodo.19034990

Bottomley, J. S., Smigelsky, M. A., Floyd, R. G., & Neimeyer, R. A. (2019). Closeness and conflict with the deceased: Exploring the factor structure of the Quality of Relationships Inventory in a bereaved sample. *Death Studies*, 43(4), 243–252.

Bowlby, J. (1969). *Attachment and Loss, Vol. 1: Attachment*. Basic Books.

Esch, T., & Stefano, G. B. (2005). The neurobiology of love. *Neuroendocrinology Letters*, 26(3), 175–192.

Field, N. P. (2006). Unresolved grief and continuing bonds: An attachment perspective. *Death Studies*, 30(8), 739–756.

Fischer, A., Halperin, E., Canetti, D., & Jasini, A. (2018). Why we hate. *Emotion Review*, 10(4), 309–320.

Fisher, H. E., Aron, A., & Brown, L. L. (2005). Romantic love: An fMRI study of a neural mechanism for mate choice. *Journal of Comparative Neurology*, 493(1), 58–62.

Gottman, J. M. (1994). *What Predicts Divorce? The Relationship Between Marital Processes and Marital Outcomes*. Lawrence Erlbaum Associates.

Heshmati, R., Zemestani, M., & Vujanovic, A. A. (2017). The role of contempt in intimate partner relationships: A brief review. *Journal of Aggression, Maltreatment & Trauma*, 26(4), 375–390.

Hewson, T., Shepperd, A., Midgley, N., & Sherr, L. (2024). Continuing bonds and grief: A systematic review. *Death Studies*, 48(3), 213–228.

Martínez, C., Bos, M. W., & Fischer, A. H. (2022). What makes hate unique? Distinctive features of hatred compared to anger and dislike. *Cognition and Emotion*, 36(8), 1535–1549.

Mikulincer, M., & Erev, I. (1991). Attachment style and the structure of romantic love. *British Journal of Social Psychology*, 30(4), 273–291.

Pretus, C., Van Bavel, J. J., & Hamlin, J. K. (2023). When does dislike become hate? The role of perceived moral violations. *Journal of Experimental Social Psychology*, 107, 104488.

Riggs, S. A. (2010). Childhood emotional abuse and the attachment system across the life cycle: What theory and research tell us. *Journal of Aggression, Maltreatment & Trauma*, 19(1), 5–51.

Riggleman, B. (2026a). Access-weighted memory decay and reconsolidation in a persistent embodied agent. Zenodo. https://doi.org/10.5281/zenodo.19122520

Riggleman, B. (2026b). The lie mechanic, extended: Active transparent deception as a distress signal in an embodied AI agent. Zenodo. https://doi.org/10.5281/zenodo.19058277

Riggleman, B. (2026d). Affective memory consolidation in a persistent embodied agent: Nightmare formation, trauma encoding, and therapeutic reconsolidation. Zenodo. https://doi.org/10.5281/zenodo.19058780

Riggleman, B. (2026h). Bias self-detection and constrained personality evolution in a persistent embodied agent. Zenodo.

Riggleman, B. (2026l). Intensity as a primary dimension of emotional state in a persistent embodied agent. Zenodo. https://doi.org/10.5281/zenodo.19159356

Riggleman, B. (2026m). The emotional geometry of a persistent agent. Zenodo. https://doi.org/10.5281/zenodo.19159429

Riggleman, B. (2026n). Addiction as frozen reconsolidation on the positive valence axis: Rock bottom as competitive co-encoding in a persistent embodied agent. Zenodo.

Root, B. L., & Exline, J. J. (2014). The role of continuing bonds in coping with grief: Overview and future directions. *Death Studies*, 38(1), 1–11.

Shapiro, A. F., Gottman, J. M., & Fink, B. C. (2015). Short-term change in couples' conflict following a transition to parenthood intervention. *Couple and Family Psychology*, 4(4), 239–251.

Simpson, J. A., & Rholes, W. S. (2017). Adult attachment, stress, and romantic relationships. *Current Opinion in Psychology*, 13, 19–24.

Smigelsky, M. A., Bottomley, J. S., Relyea, G., & Neimeyer, R. A. (2020). Investigating the relationship between grief and subjective closeness to the deceased. *Death Studies*, 44(3), 156–163.

Stroebe, M., Schut, H., & Boerner, K. (2010). Continuing bonds in adaptation to bereavement: Toward theoretical integration. *Clinical Psychology Review*, 30(2), 259–268.

Turner, H. A., Vanderminden, J., Finkelhor, D., & Hamby, S. (2019). Child neglect and the broader context of child victimization. *Child Maltreatment*, 24(3), 265–274.